

coalescent_simulation_report

March 22, 2016

1 Analysis of the Coalescent Simulation

```
In [1]: library(ggplot2)
        library(plyr)
        library(reshape2)
```

1.1 Dataset

This is the summary of Spearman's ρ over 10 replicates of the “coalescent” experiment

```
In [2]: stats = read.csv("overall.csv")
        stats$rep = as.factor(sort(rep(1:10, times=28)))
```

```
In [3]: summary(stats)
```

```
Out[3]:      coverage      measure      scale      spearman      rep
      Min.   : 0.50      ip :140      Min.   :0.001      Min.   :0.3145      1       : 28
      1st Qu.: 4.00      wip:140     1st Qu.:0.010     1st Qu.:0.7894      2       : 28
      Median :22.50                      Median :0.010     Median :0.8641      3       : 28
      Mean   :19.46                      Mean   :0.019     Mean   :0.8353      4       : 28
      3rd Qu.:30.00                      3rd Qu.:0.010     3rd Qu.:0.9159      5       : 28
      Max.   :50.00                      Max.   :0.100     Max.   :0.9727      6       : 28
                                   (Other):112
```

We compare average genome coverage and the scale of variation against accuracy (i.e. Spearman's ρ) (over the 10 reps).

We compare the effects of coverage and scale independently.


1.2 Coverage vs ρ

A series of average coverages was run at the scale of 0.01 (i.e. an average of 1 variant in 100 bases across all pairwise comparisons of samples)

```
In [4]: coverage = stats[stats$scale==0.01, ]
        coverage$scale = NULL
        summary(coverage)
```

```
Out[4]:      coverage      measure      spearman      rep
      Min.   : 0.50      ip :90      Min.   :0.3145      1       :18
      1st Qu.: 2.00      wip:90     1st Qu.:0.7592      2       :18
      Median : 8.00                      Median :0.8555      3       :18
      Mean   :13.61                      Mean   :0.8118      4       :18
      3rd Qu.:15.00                      3rd Qu.:0.9029      5       :18
      Max.   :50.00                      Max.   :0.9658      6       :18
                                   (Other):72
```

```
In [5]: ggplot(coverage, aes(x=coverage, y=spearman, linetype=measure, color=rep)) +
  geom_line(aes(linetype=measure, color=rep),size=1.5) +
  xlab('Genome Coverage') +
  ylab(expression(paste("Spearman's ", rho))) +
  #scale_x_log10()+
  theme_bw()
```



Here we summarise the replicates to averages \pm SD. Note that we exclude replicate 8 as it is an outlier for both IP and WIP metrics (see above).

```
In [6]: csumm = ddply(coverage, .(coverage, measure), summarise,
  spearman_m=mean(spearman),
  spearman_sd=sd(spearman))
summary(csumm)
```


```
Out[6]:   coverage   measure  spearman_m  spearman_sd
      Min.    : 0.50    ip : 9      Min.    :0.6385   Min.    :0.05102
```

1st Qu.: 2.00	wip:9	1st Qu.:0.7635	1st Qu.:0.07703
Median : 8.00		Median :0.8496	Median :0.09837
Mean :13.61		Mean :0.8118	Mean :0.09764
3rd Qu.:15.00		3rd Qu.:0.8772	3rd Qu.:0.10476
Max. :50.00		Max. :0.8943	Max. :0.15706

You can see below that WIP marginally outperforms IP, at low coverage. Above about 20x, I would say that WIP and IP have equivalent performance.

The ribbon is 1 SD, so there is certainly no significant difference.

```
In [7]: ggplot(csumm, aes(x=coverage, y=spearman_m, ymin=spearman_m-spearman_sd, ymax=spearman_m+spearman_sd)) +
  geom_line(aes(linetype=measure)) +
  geom_ribbon(aes(fill=measure), alpha=0.2) +
  xlab('Genome Coverage') +
  ylab(expression(paste("Spearman's ", rho))) +
  #scale_x_log10()+
  theme_bw()
```




The difference between WIP and IP is calculated here

```
In [8]: cdiff = dcast(coverage, coverage * rep ~ measure, value.var="spearman")
#cdiff = ddply(cdiff, .(coverage, rep), summarise, spearman_d=wip - ip)
cdiff = ddply(cdiff, .(coverage), summarise, diff_m=mean(wip - ip), diff_sd=sd(wip - ip))

summary(cdiff)

Out[8]:      coverage      diff_m      diff_sd
Min.   : 0.50   Min.   :-0.004446   Min.   :0.02145
1st Qu.: 2.00   1st Qu.: 0.009273   1st Qu.:0.02297
Median : 8.00   Median : 0.017816   Median :0.02738
Mean   :13.61   Mean   : 0.018810   Mean   :0.02696
3rd Qu.:15.00   3rd Qu.: 0.034508   3rd Qu.:0.03065
Max.   :50.00   Max.   : 0.042950   Max.   :0.03547

In [9]: ggplot(cdiff, aes(x=coverage, y=diff_m, ymin=diff_m-diff_sd, ymax=diff_m+diff_sd)) +
  geom_line() +
  geom_ribbon(alpha=0.4) +
  xlab('Genome Coverage') +
  ylab(expression(paste("Spearman's ", rho))) +
  #scale_x_log10()+
  theme_bw()
```



1.3 Scale vs ρ

Like coverage, we investigate the effect of variation at a constant coverage, in this case 30x. I also convert the scale into its inverse, as this is how some people prefer to think of it (i.e. one variant in X bases, as opposed to 0.0x variants per base on average. Each to their own...)


```
In [10]: scale = stats[stats$coverage==30, ]
         scale$scale = 1/scale$scale
```

```
In [11]: ssumm = ddpdy(scale, .(scale, measure), summarise,
                        spearman_m=mean(spearman),
                        spearman_sd=sd(spearman))
         summary(csumm)
```

```
Out[11]:   coverage   measure  spearman_m  spearman_sd
         Min.    : 0.50    ip :9    Min.    :0.6385    Min.    :0.05102
```

1st Qu.: 2.00	wip:9	1st Qu.:0.7635	1st Qu.:0.07703
Median : 8.00		Median :0.8496	Median :0.09837
Mean :13.61		Mean :0.8118	Mean :0.09764
3rd Qu.:15.00		3rd Qu.:0.8772	3rd Qu.:0.10476
Max. :50.00		Max. :0.8943	Max. :0.15706

```
In [12]: ggplot(ssumm, aes(x=scale, y=spearman_m, ymin=spearman_m-spearman_sd, ymax=spearman_m+spearman_sd)) +
  geom_line(aes(linetype=measure)) +
  geom_ribbon(aes(fill=measure), alpha=0.2) +
  xlab('Divergence (1 variant per x bases)') +
  ylab(expression(paste("Spearman's ", rho))) +
  #scale_x_log10() +
  scale_x_reverse() +
  theme_bw()
```



1.4 Conclusions

- I think there might be an issue with the way I normalise trees. I think that we are probably at a higher level of divergence than I expect if we use the mean. I will do a run with a couple of reps using the maximum distance set to 1.0, i.e. that the entire tree scale is 0.5 (from root to tip, and then back again =1.0).
- I'd like to re-do the coverage sweep at a scale of 0.005 or 0.002 or even 0.001. I think that this might be more inline with our rice experiment. My take home from this is that WIP is only important when your signal:noise ratio is low, like when you have a small amount of variation. Otherwise, they are equivalent (neither is significantly worse on average). Norman, can you comment?

In []: